

Parallel Computing Mid-Term: *k-means*

Giovanni Bindi

Università degli Studi di Firenze

24 February 2021



Introduction: *k-means*

K-means is a simple clustering algorithm developed during the 1950s [1]: given a set of observations $\mathcal{S} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ with $\mathbf{x} \in \mathbb{R}^D$ the algorithm aims to find $k \leq N$ sets $\mathcal{C} = \{\mathcal{S}_1, \dots, \mathcal{S}_k\}$ as to:


$$\mathcal{C}^* = \arg \min_{\mathcal{C}} \sum_{i=1}^k \sum_{\mathbf{x}_i \in \mathcal{S}_i} \|\mathbf{x}_i - \boldsymbol{\mu}_i\|_2^2 \quad (1)$$

where $\boldsymbol{\mu}_i$ is the mean of the points in \mathcal{S}_i :

$$\boldsymbol{\mu}_i = \frac{1}{|\mathcal{S}_i|} \sum_{\mathbf{x}_j \in \mathcal{S}_i} \mathbf{x}_j \quad (2)$$

with the conditions

$$\begin{cases} \bigcup_{i=1}^k \mathcal{S}_i = \mathcal{S} \\ \mathcal{S}_i \cap \mathcal{S}_j = \emptyset, & i \neq j \end{cases} \quad (3)$$



Algorithm (Lloyd's version)

Drawback

Problem 1 is NP-hard [2] with a time complexity of $O(N^{ND+1})$ [3]

→ Need for an heuristic.

Lloyd's algorithm: after initial sampling phase $\mathcal{C}^{(1)} = \{\mathcal{S}_k^{(1)}, \dots, \mathcal{S}_k^{(1)}\}$ for $t = 1, \dots, l$ iterations repeat two phases:

- 1 **Assignment** step: where each data point is assigned to the cluster with the nearest mean, i.e. building $\mathcal{S}_i^{(t)}$:

$$\mathcal{S}_i^{(t)} = \{\mathbf{x}_p : \|\mathbf{x}_p - \boldsymbol{\mu}_i^{(t)}\|^2 \leq \|\mathbf{x}_p - \boldsymbol{\mu}_j^{(t)}\|_2^2\} \quad (4)$$

for all $\mathbf{x}_p \in \mathcal{S}$ and for all $j \neq i$.

- 2 **Update** step: where the centroids are recomputed, based on the previous step calculations:

$$\boldsymbol{\mu}_i^{(t+1)} = \frac{1}{|\mathcal{S}_i^{(t)}|} \sum_{\mathbf{x}_j \in \mathcal{S}_i^{(t)}} \mathbf{x}_j \quad (5)$$

Pseudo-code: Sequential Version

Input: $\mathcal{S} \subset \mathbb{R}^D$, $k, l \in \mathbb{N}$. Optional: $\epsilon \in \mathbb{R}$.

Output: Clusters \mathcal{C} with centroids $\{\mu_1, \dots, \mu_k\}$

```
1:  $\mu^{(1)} = \{\mu_1^{(1)}, \dots, \mu_k^{(1)}\} \leftarrow \text{sample}(\mathcal{S})$ 
2: for  $t = 1, \dots, l$  do
3:    $\tau = \{0, \dots, 0\}$ 
4:    $\eta = \{0, \dots, 0\}$ 
5:   for  $i = 1, \dots, N$  do
6:      $c = \text{closest\_centroid}(\mathbf{x}_i, \mu^{(t)})$ 
7:      $\tau_c = \tau_c + \mathbf{x}_i$ 
8:      $\eta_c = \eta_c + 1$ 
9:   end for
10:  for  $j = 1, \dots, k$  do
11:     $\mu_j^{(t+1)} = \tau_j / \eta_j$ 
12:  end for
13:  if  $\|\mu^{(t+1)} - \mu^{(t)}\|_2^2 \leq \epsilon$  then
14:    return  $\mu^{(t+1)}$ 
15:  end if
16: end for
17: return  $\mu^{(l)}$ 
```

Pseudo-code: Parallel Version

Input: $S \subset \mathbb{R}^D$, $k, l \in \mathbb{N}$. $T \in \mathbb{N}$. Optional: $\epsilon \in \mathbb{R}$.

Output: Clusters \mathcal{C} with centroids $\{\mu_1, \dots, \mu_k\}$

```

1:  $\mu^{(1)} = \{\mu_1^{(1)}, \dots, \mu_k^{(1)}\} \leftarrow \text{sample}(S)$ 
2:  $\mathcal{P} \leftarrow \text{init\_pool}(T)$ 
3: for  $t = 1, \dots, l$  do
4:    $\tau = \{0, \dots, 0\}$ 
5:    $\eta = \{0, \dots, 0\}$ 
6:   for each thread  $p \in \mathcal{P}$  do
7:      $(s, e) \leftarrow \text{compute\_start\_end}(p, N)$ 
8:     for  $i = s, \dots, e$  do
9:        $c = \text{closest\_centroid}(\mathbf{x}_i, \mu^{(t)})$ 
10:       $\tau_c = \tau_c + \mathbf{x}_i$ 
11:       $\eta_c = \eta_c + 1$ 
12:    end for
13:  end for
14:   $\text{wait\_all\_threads}()$ 
15:  for  $j = 1, \dots, k$  do
16:     $\mu_j^{(t+1)} = \tau_j / \eta_j$ 
17:  end for
18:  if  $\|\mu^{(t+1)} - \mu^{(t)}\|_2^2 \leq \epsilon$  then
19:    return  $\mu^{(t+1)}$ 
20:  end if
21: end for
22: return  $\mu^{(l)}$ 

```

Thread Pool

```

class thread_pool {
    private:
        std::vector<std::thread> workers;
        std::deque<std::function<void()>> tasks;
        std::mutex mtx;
        std::condition_variable task_done;
        std::condition_variable all_done;

        unsigned int running_tasks;
        bool stop;

        void run_task() {
            while (true) {
                std::unique_lock<std::mutex> lock(mtx);
                task_done.wait(lock, [this]() { return stop || !tasks.empty(); });
                if (!tasks.empty()) {
                    ++running_tasks;
                    auto func = tasks.front();
                    tasks.pop_front();
                    lock.unlock();
                    func();
                    lock.lock();
                    --running_tasks;
                    all_done.notify_one();
                }
                else if (stop)
                    break;
            }
            return;
        }
    ...
}

```

Thread Pool - 2

```

...
public:

    thread_pool(const unsigned int T = std::thread::hardware_concurrency()) :
        running_tasks(),
        stop() {
            for (unsigned int i = 0; i < T; ++i)
                workers.emplace_back(std::bind(&thread_pool::run_task, this));
        }

    ~thread_pool() {
        std::unique_lock<std::mutex> lock(mtx);
        stop = true;
        task_done.notify_all();
        lock.unlock();
        for (auto& w : workers)
            w.join();
    }

    template<class F>
    void enqueue(F&& f) {
        std::unique_lock<std::mutex> lock(mtx);
        tasks.emplace_back(std::forward<F>(f));
        task_done.notify_one();
        return;
    }

    void wait_all_threads() {
        std::unique_lock<std::mutex> lock(mtx);
        all_done.wait(lock, [this]() { return tasks.empty() && (running_tasks == 0); });
        return;
    }
};

```

Experiments

- ▶ 4C / 8T CPU (Intel i7-8550U).
- ▶ No ϵ -stop $\rightarrow I = 1000$ iterations.
- ▶ Ten runs for each experiment:


$$\bar{t}_A = \frac{1}{10} \sum_{i=1}^{10} t_A^{(i)} \quad (6)$$

- ▶ Speedup of B over A :

$$S = \bar{t}_A / \bar{t}_B \quad (7)$$

- ▶ Variance:

$$\sigma^2 \approx S^2 \left[\frac{\sigma_A^2}{\bar{t}_A^2} + \frac{\sigma_B^2}{\bar{t}_B^2} \right] \quad (8)$$



- ▶ Four dataset sizes $N = \{10000, 50000, 100000, 500000\}$.
- ▶ Two dimensionalities $D = \{2, 3\}$.
- ▶ Three centroids.

General Behaviour

Algorithm time complexity: $O(NDIk)$

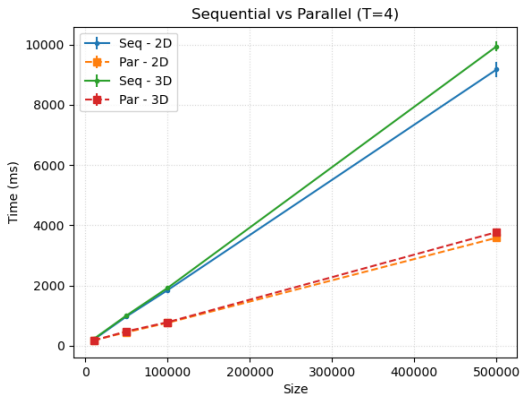


Figure 1: Execution time (ms): Sequential vs Parallel (with $T = 4$)

Against Sequential

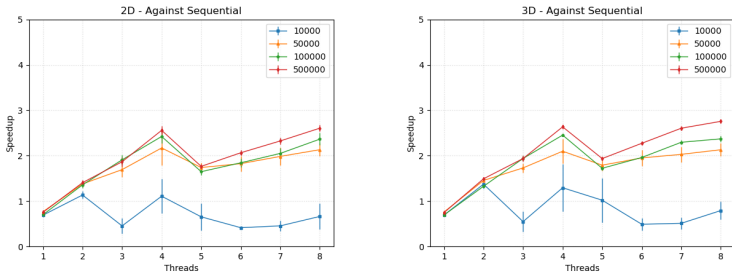


Figure 2: Speedup: Parallel vs Sequential for 2D (left) and 3D (right) datasets

- ▶ **Maximum** speedup: $S^{max} = 2.76$ for $N = 5 \cdot 10^5$, $D = 3$ and $T = 8$.
- ▶ **Minimum** speedup: $S^{min} = 0.42$ for $N = 1 \cdot 10^4$, $D = 2$ and $T = 6$.

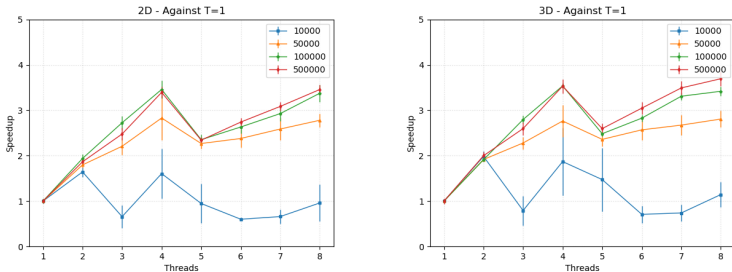
Against $T = 1$ 

Figure 3: Speedup: Parallel vs $T = 1$ for 2D (left) and 3D (right) datasets

- ▶ **Maximum** speedup: $S^{max} = 3.69$ for $N = 5 \cdot 10^5$, $D = 3$ and $T = 8$.
- ▶ **Minimum** speedup: $S^{min} = 0.60$ for $N = 1 \cdot 10^4$, $D = 2$ and $T = 6$

Conclusions

Two implementations of *k-means*: sequential and parallel.

Pros:

- ▶ Parallelization method: thread pool & task queue.
- ▶ Header-only C++17 libraries, with simple test cases.
- ▶ Passes code checkers: *Cppcheck* and *Valgrind*.
- ▶ Obtained a perceivable speedup.

Cons:

- ▶ Speedup is small.

Code

<https://github.com/w00zie/kmeans>

References I

- [1] S. Lloyd, “Least squares quantization in pcm,” *IEEE Transactions on Information Theory*, vol. 28, no. 2, pp. 129–137, 1982. DOI: [10.1109/TIT.1982.1056489](https://doi.org/10.1109/TIT.1982.1056489).
- [2] D. Aloise, A. Deshpande, P. Hansen, and P. Popat, “Np-hardness of euclidean sum-of-squares clustering,” *Mach. Learn.*, vol. 75, no. 2, pp. 245–248, May 2009, ISSN: 0885-6125. DOI: [10.1007/s10994-009-5103-0](https://doi.org/10.1007/s10994-009-5103-0). [Online]. Available: <https://doi.org/10.1007/s10994-009-5103-0>.
- [3] M. Inaba, N. Katoh, and H. Imai, “Applications of weighted voronoi diagrams and randomization to variance-based k-clustering: (extended abstract),” in *Proceedings of the Tenth Annual Symposium on Computational Geometry*, ser. SCG '94, Stony Brook, New York, USA: Association for Computing Machinery, 1994, pp. 332–339, ISBN: 0897916484. DOI: [10.1145/177424.178042](https://doi.org/10.1145/177424.178042). [Online]. Available: <https://doi.org/10.1145/177424.178042>.

Code

Data Structures:

```
template <typename T, const size_t D>
using vec = std::array<T, D>;

template <typename T, const size_t N, const size_t D>
using mat = std::array<vec<T, D>, N>;

template <typename T, const size_t D>
using hash = tsl::robin_map<size_t, vec<T, D>>;

using counter = tsl::robin_map<size_t, size_t>;
```

Distance calculation:

```
template<typename T, const size_t D>
double calc_distance(const vec<T, D>& p, const vec<T, D>& q) {
    double sum = 0.0;
    for (size_t i = 0; i < D; ++i)
        sum += ((p[i] - q[i]) * (p[i] - q[i]));
    return sum;
}

template<typename T, const size_t M, const size_t D>
size_t get_closest_centroid(const vec<T, D>& point, hash<T, D>& centroids) {
    vec<double, M> distances;
    for (size_t i = 0; i < M; ++i)
        distances[i] = calc_distance(point, centroids[i]);
    auto min_dist = std::min_element(distances.begin(), distances.end());
    return std::distance(distances.begin(), min_dist);
}
```

Code - Sequential

Core of sequential code:

```
template <typename T, const size_t N, const size_t M, const size_t D>
hash<T, D> assign_and_update(const mat<T, N, D>& data, hash<T, D>& centroids) {
    hash<T, D> new_centroids;
    counter counts;
    new_centroids.reserve(M);
    counts.reserve(M);
    // Assignment step
    for (size_t i = 0; i < N; ++i) {
        size_t closest_centroid = get_closest_centroid<T, M, D>(data[i], centroids);
        new_centroids[closest_centroid] = new_centroids[closest_centroid] + data[i];
        counts[closest_centroid] += 1;
    }
    // Update step
    for (size_t i = 0; i < M; ++i)
        new_centroids[i] = new_centroids[i] / counts[i];
    return new_centroids;
}

template <typename T, const size_t N, const size_t M, const size_t D>
hash<T, D> kmeans(const mat<T, N, D>& data, hash<T, D>& centroids, const size_t niter) {
    for (size_t i = 0; i < niter; ++i)
        centroids = assign_and_update<T, N, M, D>(data, centroids);
    return centroids;
}
```

Code - Parallel

Core of parallel code - 1:

```

std::mutex mtx;

template <typename T, const size_t N, const size_t M, const size_t D>
void assign_subset(const mat<T, N, D>& data, hash<T, D>& centroids,
                  hash<T, D>& future_centroids, counter& global_counts,
                  const size_t start, const size_t end) {

    hash<T, D> partial_centroids;
    counter partial_counts;
    partial_centroids.reserve(M);
    partial_counts.reserve(M);
    // Partial assignment
    for (size_t i = start; i < end; ++i) {
        size_t closest_centroid = get_closest_centroid<T, M, D>(data[i], centroids);
        partial_centroids[closest_centroid] = partial_centroids[closest_centroid] + data[i];
        partial_counts[closest_centroid] += 1;
    }
    std::lock_guard<std::mutex> lock(mtx);
    for (const auto& c : partial_centroids) {
        future_centroids[c.first] = future_centroids[c.first] + c.second;
        global_counts[c.first] += partial_counts[c.first];
    }
    return;
}

```


Code - Parallel 2

Core of parallel code - 2:

```

template <typename T, const size_t N, const size_t M, const size_t D>
hash<T, D> kmeans(const mat<T, N, D>& data, hash<T, D>& centroids,
                 const size_t niter, const size_t nthreads) {
    kmeans::par::tp::thread_pool pool(nthreads);
    const size_t step = (N + nthreads - 1) / nthreads;
    for (size_t i = 0; i < niter; ++i) {
        hash<T, D> future_centroids;
        counter global_counter;
        future_centroids.reserve(M);
        global_counter.reserve(M);
        for (size_t t = 0; t < N; t += step) {
            size_t real_step = (t + step > N) ? N - t : step;
            pool.enqueue([&, t, real_step]() {
                assign_subset<T, N, M, D>(data, centroids, future_centroids,
                                          global_counter, t, t+real_step);
            });
        }
        pool.wait_all_threads();
        // Update step
        for (const auto& fc : future_centroids)
            centroids[fc.first] = fc.second / global_counter[fc.first];
    }
    return centroids;
}

```